File   Edit   View   Insert   Cell   Kernel   Help   Not Trusted | Python 3 ○

N-grams Corpus preprocessing

The input corpus in this week's assignment is a continuous text that needs some preprocessing so that you can start calculating the n-gram probabilities.

Some common preprocessing steps for the language models include:

- lowercasing the text
- remove special characters
- split text to list of sentences
- split sentence into list words

Can you note the similarities and differences among the preprocessing steps shown during the Course 1 of this specialization?

```python
In [1]: import nltk              # NLP toolkit
        import re                # Library for Regular expression operations

        nltk.download('punkt')   # Download the Punkt sentence tokenizer
```

```
[nltk_data] Downloading package punkt to /home/jovyan/nltk_data...
[nltk_data]   Unzipping tokenizers/punkt.zip.
```

Out[1]: True

Lowercase

Words at the beginning of a sentence and names start with a capital letter. However, when counting words, you want to treat them the same as if they appeared in the middle of a sentence.

You can do that by converting the text to lowercase using [str.lowercase] (https://docs.python.org/3/library/stdtypes.html?highlight=split#str.lower).

```python
In [2]: # change the corpus to lowercase
        corpus = "Learning% makes 'me' happy. I am happy be-cause I am learning! :)"
        corpus = corpus.lower()

        # note that word "Learning" will now be the same regardless of its position in the sentence
        print(corpus)
```

```
learning% makes 'me' happy. i am happy be-cause i am learning! :)
```

Remove special charactes

Some of the characters may need to be removed from the corpus before we start processing the text to find n-grams.

Often, the special characters such as double quotes '"' or dash '-' are removed, and the interpunction such as full stop '.' or question mark '?' are left in the corpus.

```python
In [3]: # remove special characters
        corpus = "learning% makes 'me' happy. i am happy be-cause i am learning! :)"
        corpus = re.sub(r"[^a-zA-Z0-9.?! ]+", "", corpus)
        print(corpus)
```

```
learning makes me happy. i am happy because i am learning!
```

Note that this process gets rid of the happy face made with punctuations :). Remember that for sentiment analysis, this emoticon was very important. However, we will not consider it here.

Text splitting

In the assignment, the sentences in the corpus are separated by a special delimiter \n. You will need to split the corpus into an array of sentences using this delimiter. One way to do that is by using the str.split method.

The following examples illustrate how to use this method. The code shows:

- how to split a string containing a date into an array of date parts
- how to split a string with time into an array containing hours, minutes and seconds

Also, note what happens if there are several back-to-back delimiters like between "May" and "9".

```python
In [4]: # split text by a delimiter to array
        input_date="Sat May  9 07:33:35 CEST 2020"

        # get the date parts in array
        date_parts = input_date.split(" ")
        print(f"date parts = {date_parts}")

        #get the time parts in array
        time_parts = date_parts[4].split(":")
        print(f"time parts = {time_parts}")
```

```
date parts = ['Sat', 'May', '', '9', '07:33:35', 'CEST', '2020']
time parts = ['07', '33', '35']
```

This text splitting is more complicated than the tokenization process used for sentiment analysis.

Sentence tokenizing

Once you have a list of sentences, the next step is to split each sentence into a list of words.

This process could be done in several ways, even using the str.split method described above, but we will use the NLTK library nltk to help us with that.

In the code assignment, you will use the method word_tokenize to split your sentence into a list of words. Let us try the method in an example.

```
In [5]: # tokenize the sentence into an array of words

sentence = 'i am happy because i am learning.'
tokenized_sentence = nltk.word_tokenize(sentence)
print(f'{sentence} -> {tokenized_sentence}')
```

```
i am happy because i am learning. -> ['i', 'am', 'happy', 'because', 'i', 'am', 'learning', '.']
```

Now that the sentence is tokenized, you can work with each word in the sentence separately. This will be useful later when creating and counting N-grams. In the following code example, you will see how to find the length of each word.

```
In [6]: # find length of each word in the tokenized sentence
sentence = ['i', 'am', 'happy', 'because', 'i', 'am', 'learning', '.']
word_lengths = [(word, len(word)) for word in sentence] # Create a list with the word lengths using a list comprehension
print(f' Lengths of the words: \n{word_lengths}')
```

```
 Lengths of the words:
[('i', 1), ('am', 2), ('happy', 5), ('because', 7), ('i', 1), ('am', 2), ('learning', 8), ('.', 1)]
```

The previous result produces a list of pairs. This is not equivalent to a dictionary.

N-grams

Sentence to n-gram

The next step is to build n-grams from the tokenized sentences.

A sliding window of size n-words can generate the n-grams. The window scans the list of words starting at the sentence beginning, moving by a step of one word until it reaches the end of the sentence.

Here is an example method that prints all trigrams in the given sentence.

```
In [7]: def sentence_to_trigram(tokenized_sentence):
    """
    Prints all trigrams in the given tokenized sentence.

    Args:
        tokenized_sentence: The words list.

    Returns:
        No output
    """
    # note that the last position of i is 3rd to the end
    for i in range(len(tokenized_sentence) - 3 + 1):
        # the sliding window starts at position i and contains 3 words
        trigram = tokenized_sentence[i : i + 3]
        print(trigram)

tokenized_sentence = ['i', 'am', 'happy', 'because', 'i', 'am', 'learning', '.']

print(f'List all trigrams of sentence: {tokenized_sentence}\n')
sentence_to_trigram(tokenized_sentence)
```

```
List all trigrams of sentence: ['i', 'am', 'happy', 'because', 'i', 'am', 'learning', '.']

['i', 'am', 'happy']
['am', 'happy', 'because']
['happy', 'because', 'i']
['because', 'i', 'am']
['i', 'am', 'learning']
['am', 'learning', '.']
```

Prefix of an n-gram

As you saw in the lecture, the n-gram probability is often calculated based on the (n-1)-gram counts. The prefix is needed in the formula to calculate the probability of an n-gram.

$$P(w_n|w_1^{n-1}) = \frac{C(w_1^n)}{C(w_1^{n-1})}$$

The following code shows how to get an (n-1)-gram prefix from n-gram on an example of getting trigram from a 4-gram.

```
In [8]: # get trigram prefix from a 4-gram
fourgram = ['i', 'am', 'happy','because']
trigram = fourgram[0:-1] # Get the elements from 0, included, up to the last element, not included.
print(trigram)
```

```
['i', 'am', 'happy']
```

Start and end of sentence word $< s >$ and $< /s >$

You could see in the lecture that we must add some special characters at the beginning and the end of each sentence:

- $<s>$ at beginning
- $</s>$ at the end

For n-grams, we must prepend n-1 of characters at the begining of the sentence.

Let us have a look at how you can implement this in code.

```python
# when working with trigrams, you need to prepend 2 <s> and append one </s>
n = 3
tokenized_sentence = ['i', 'am', 'happy', 'because', 'i', 'am', 'learning', '.']
tokenized_sentence = ["<s>"] * (n - 1) + tokenized_sentence + ["</s>"]
print(tokenized_sentence)
```

```
['<s>', '<s>', 'i', 'am', 'happy', 'because', 'i', 'am', 'learning', '.', '</s>']
```

That's all for the lab for "N-gram" lesson of week 3.